

Message-Oriented Middleware

Linh B. Ngo
March 11, 2016



Distributed Systems [1]

- Software systems (e.g., sensor systems) are distributed
 - Increasing scales
 - Across geographical and organizational boundaries
- How to scale up the communication infrastructure accordingly?

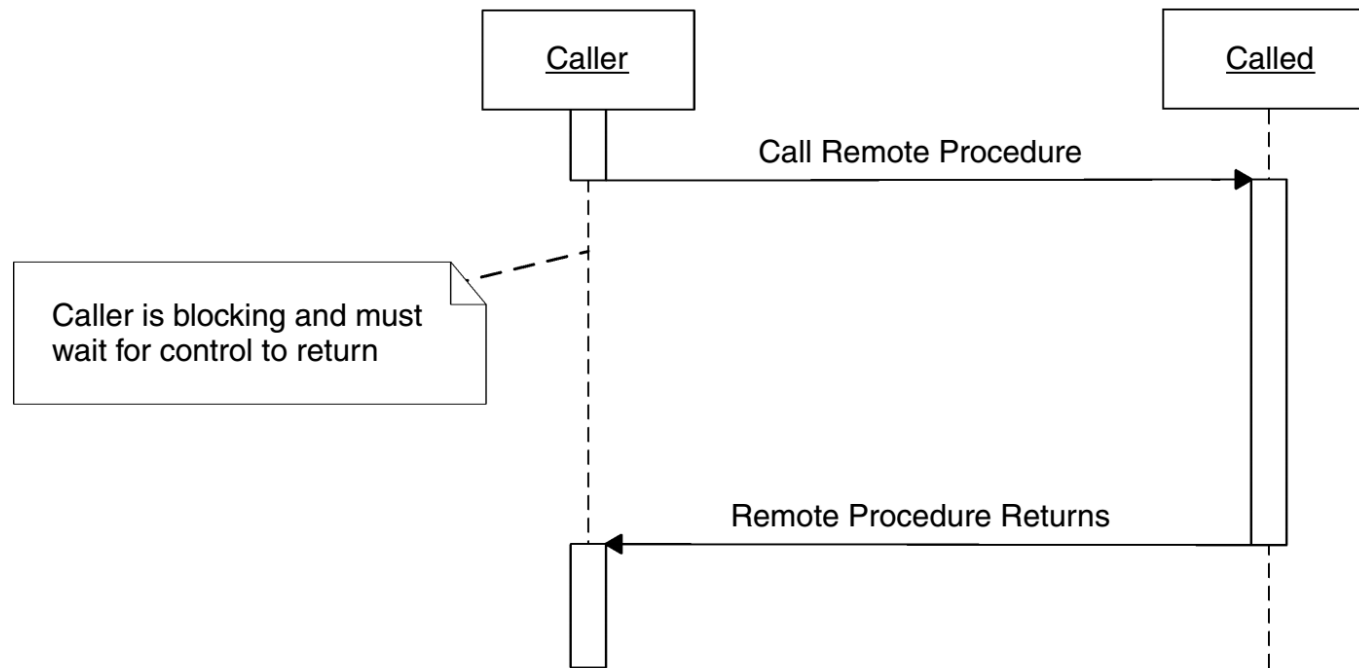
Challenges

- Complex environments
- Multiple programming languages
- Multiple hardware platforms (computing nodes and sensor platforms)
- Multiple operating systems
- Dynamic and flexible deployment
- High reliability, throughput, and resiliency
- High quality of service

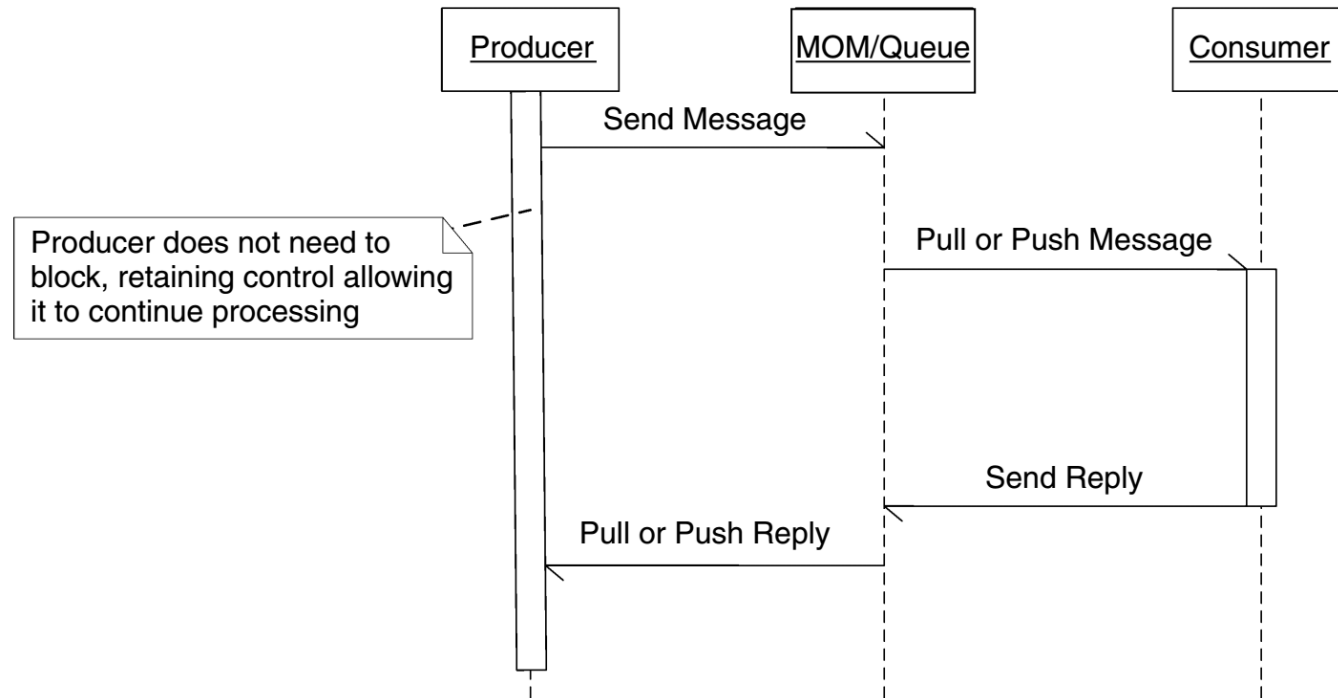
MOM: Message Oriented Middleware

- Alternative to RPC distribution mechanism
- Clean method for communication among disparate software entities
- Peer-to-peer relationship between individual clients (entities)


Synchronous




Asynchronous



Example RPC



Example MOM



RPC vs MOM: Coupling

- RPC is designed to work on object or function interfaces, resulting in tightly coupled systems
- The intermediate layer between senders and receivers in MOMs allows entities to exchange messages in a loosely coupled manner.

RPC vs MOM: Reliability

- Most RPC implementations provide little or no guaranteed reliable communication capability
- MOM relies on a store-and-forward mechanism to prevent message loss through network or system failure

RPC vs MOM: Scalability

- Due to its blocking nature, RPC can impact system performance when the participating components do not scale equally (bottleneck due to synchronization)
- As the subsystems are decoupled, MOM also decouple the performance dependency among them, resulting in the individual scaling ability.

RPC vs MOM: Availability

- The synchronous characteristics of RPC are often led to interdependency among subsystems. Consequently, the impact of individual component failures can be propagated across the entire system.
- Without a rigid coupling of subsystems (due to synchronous communication), MOM reduces failure propagation.

MOM: Message Queues

- Fundamental concept within MOM
- FIFO
- Customizable:
 - Queue Name
 - Queue Size
 - Saving Threshold
 - Message-sorting algorithm
 - ...

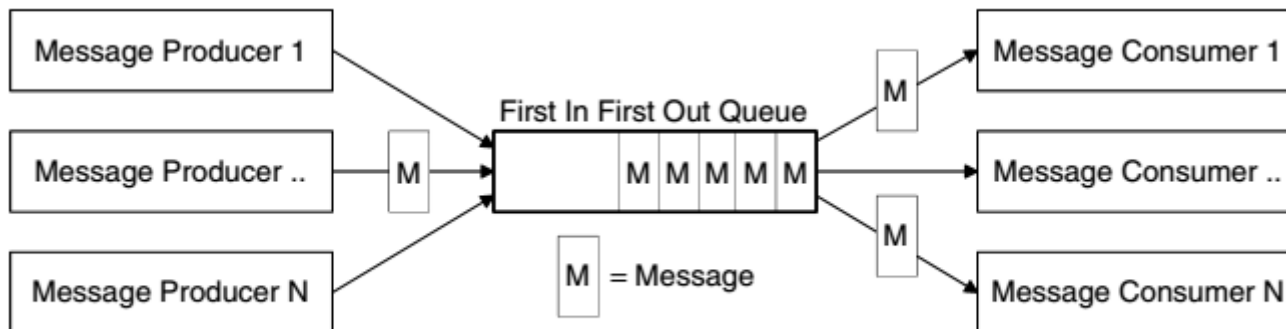

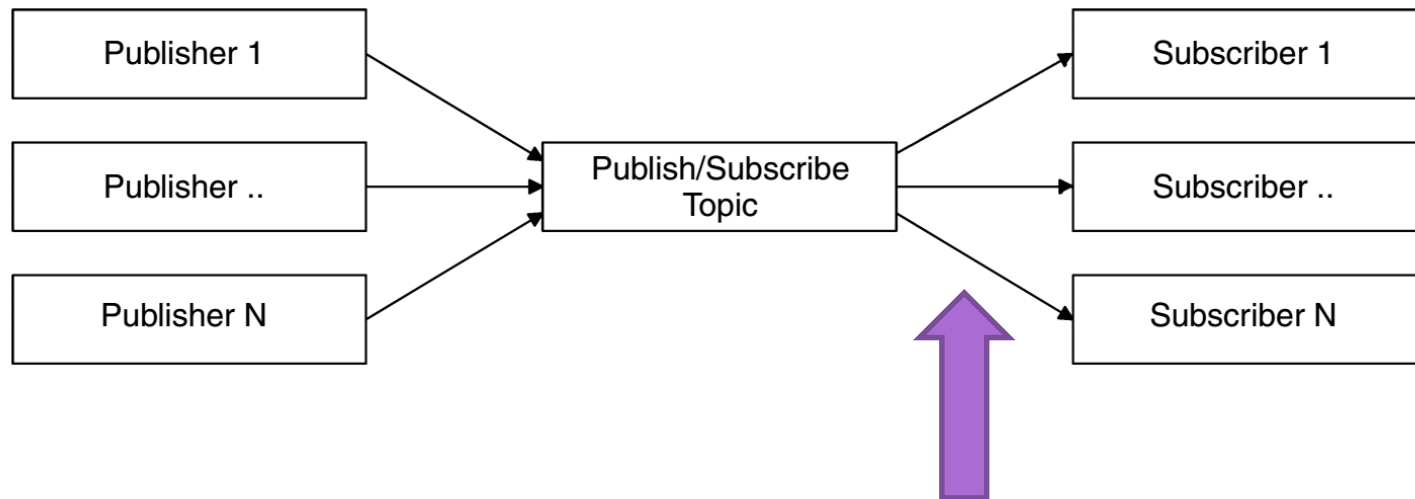


Figure 1.5 Message queue

Point-to-Point Messaging Model



Publish/Subscribe Model



Subscribers can **pull** from the Topic or have the messages **pushed** to themselves

Common MOM Standards

- Advanced Message Queueing Protocol (AMQP)
 - ISO standard
 - Application layer protocol
- MQ Telemetry Transport (MQTT)
 - ISO standard
 - Lightweight publish/subscribe messaging transport protocol on top of TCP/IP for M2M/IoT contexts
- Data Distribution Service
 - Message-oriented PubSub middleware standard interfacing with C++, C++11, C, Ada, Java, and Ruby
- Extensible Messaging and Presence Protocol (XMPP)
 - Message-oriented middleware based on XML.
 - Open standard

Open Source MOMs

- RabbitMQ
 - Apache ActiveMQ
 - ZeroMQ
 - Kafka
 - Apache Qpid
 - IronMQ
-
- No in-depth performance comparison
 - Different message/bandwidth characteristics call for different MOMs

RabbitMQ

- AMQP
- Erlang
- Broker architecture

Apache ActiveMQ

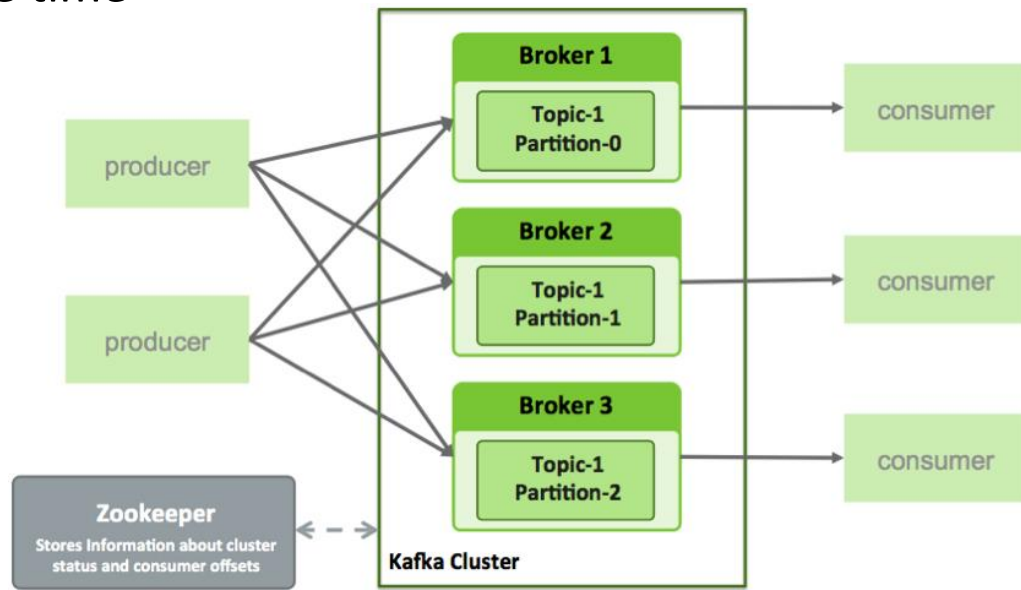
- AMQP, MQTT, Openwire, WebSockets ...
- Java

ZeroMQ

- Very lightweight
- Require manual implementation
- ZMTP

Apache Kafka [2]

- Open-source message broker project (LinkedIn)
- Unified, high-throughput, low-latency platform to handle real-time data feeds
- A topic is divided into multiple partitions, and each broker stores one or more of these partitions
- Multiple producers/consumers can publish and retrieve messages at the same time



Apache Kafka: Efficiency on Single Partition

- Simple Storage
 - Each partition of a topic correspond to a logical log
 - A logical log is a set of segment files of same size (1GB)
 - Segment files are flushed to disk after a number of messages has been published (or after a determined duration)
 - Messages are exposed to consumers after they are flushed
 - No message id: message is addressed by its offset in the log
- Efficient Transfer
 - No application-level caching
 - Sequential data access
- Stateless Broker
 - Records of data consumed are not maintained
 - Data are purged from brokers after default length of 7 days

Apache Kafka: Distributed Coordination

- Smallest unit of parallelism: partition
 - Each partition can only be consumed by one consumer at any given time
 - More partitions than consumers
- No central master node – Zookeeper instead
 - Detecting the addition and removal of brokers and consumers
 - Triggering rebalancing in consumers
 - Maintaining consumption relationship and keep track of the consumed offset of each partition



Apache Kafka: Delivery Guarantee

- At-least-once
- Most of the time is exactly-one
- Duplication happens when consumers crash and is revived.
- Messages from a single partition is guarantee to be delivered in-order

Apache Kafka: Performance Evaluation

- 2 machines: 8 2Ghz cores, 16GB RAM, 6 HDD (RAID 10)
- 1GB network
- *Note: No simultaneous producing/consuming test*
- Producer Test:
 - 10M messages, size 200 bytes
 - Batches of size 1 and 50
 - Results: 50K/s for size 1 and 400K/s for size 50
- Consumer Test:
 - 10M messages, size 200 bytes
 - Pull requests prefetch 1000 messages
 - Results: 22K/s

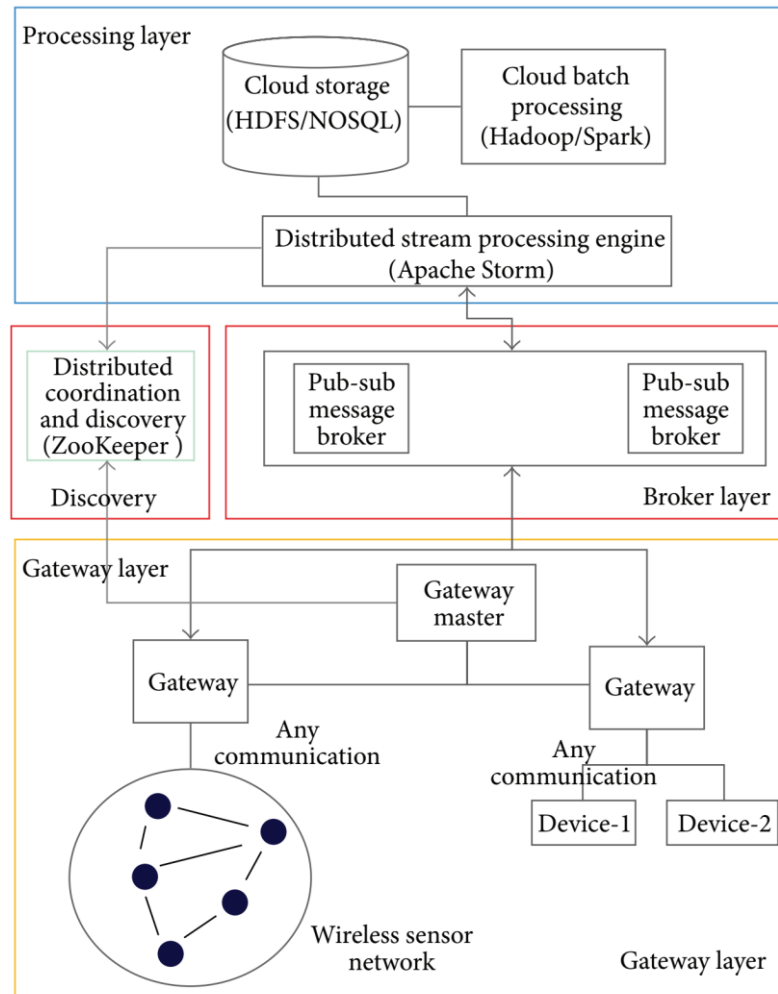
Apache Kafka: Performance Evaluation



Apache Kafka: Performance Evaluation

- More efficient storage format
- Stateless broker
- sendfile API reduces overhead

Apache Kafka: Another Evaluation



Apache Kafka: Another Evaluation

- FutureGrid platform
- 8 virtual nodes: 2VCPU, 4GBRAM, 40GB HDD
- 1 node: Storm Nimbus and Zookeeper
- 2 nodes: Gateway servers
- 3 nodes: Storm supervisors
- 2 nodes: Brokers

Apache Kafka: Another Evaluation [3]

- FutureGrid platform
- 8 virtual nodes: 2VCPU, 4GBRAM, 40GB HDD
- 1 node: Storm Nimbus and Zookeeper
- 2 nodes: Gateway servers
- 3 nodes: Storm supervisors
- 2 nodes: Brokers
- 4 applications deployed on the two gateways producing data with constant rates:
 - Up to 100 messages per second
 - Up to 1MB in message size
 - Data are relayed into Storm, and then passed back to the gateway
- Timing measurements: roundtrip

Apache Kafka: Another Evaluation

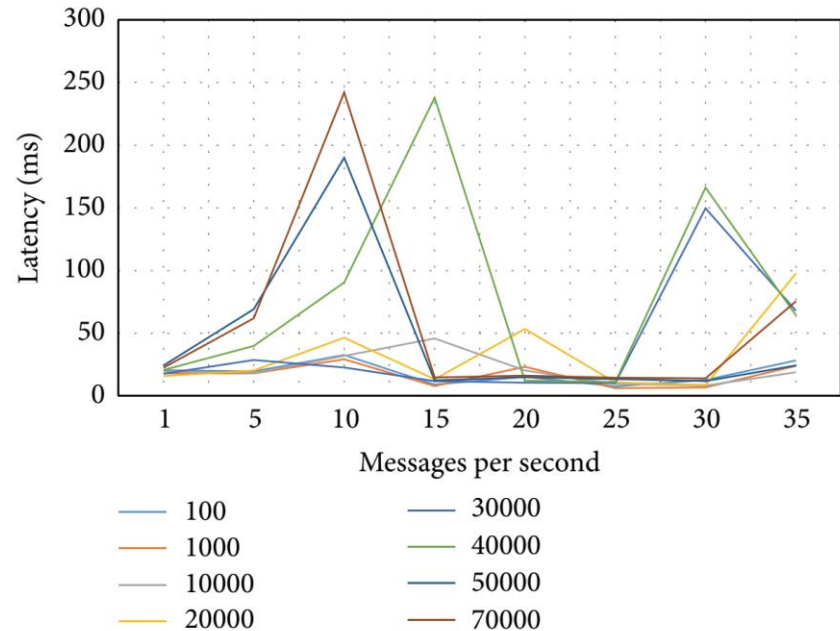
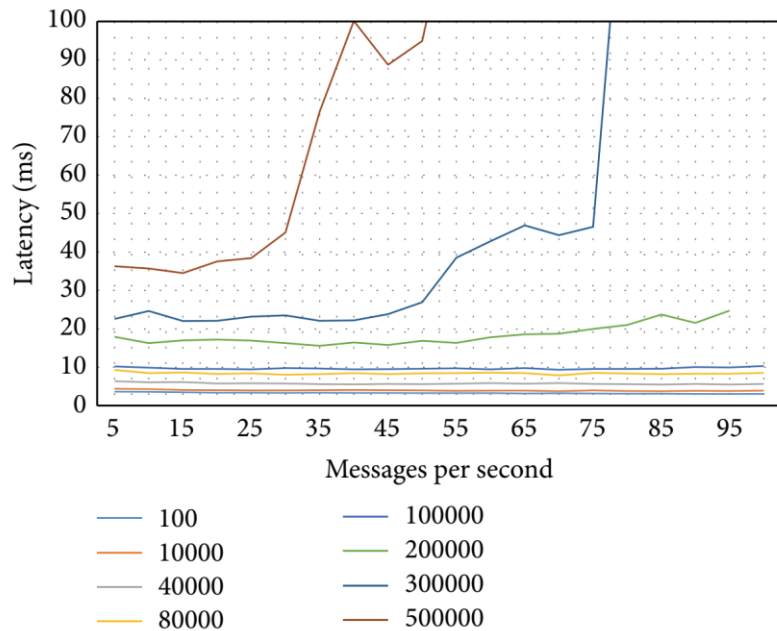


FIGURE 7: Average latency for different message sizes with RabbitMQ. The different lines are for different message sizes in bytes.

FIGURE 8: Average latency for different message sizes with Kafka. The different lines are for different message sizes in bytes.

Apache Kafka: Another Evaluation




FIGURE 10: Latency with varying number of devices, RabbitMQ. The average latency and standard deviation are shown.

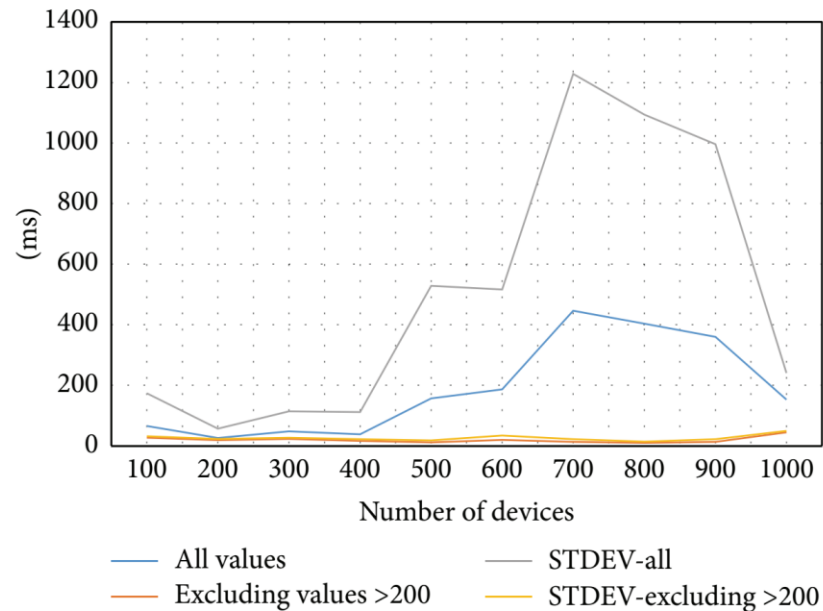


FIGURE 11: Latency with varying number of devices, Kafka. The average latency and standard deviation are shown. Also, averages calculated with omitting values over 200 are shown.

References

1. Curry, Edward. "Message-oriented middleware." *Middleware for communications* (2004): 1-28.
2. Kreps, Jay, Neha Narkhede, and Jun Rao. "Kafka: A distributed messaging system for log processing." *NetDB*, 2011
3. Kamburugamuve, Supun, Leif Christiansen, and Geoffrey Fox. "A framework for real time processing of sensor data in the cloud." *Journal of Sensors* 2015 (2015).