

# MapReduce Design Patterns

(Based on “MapReduce Design Patterns” by  
Donald Miner and Adam Shook)

Linh B. Ngo

# CONTENTS

- Motivation
- Basic Patterns
- Summarization
- Filtering
- Data Organization
- Join
- Metapatterns

# MOTIVATION

- MapReduce is designed as a framework
  - The solution has to fit into the framework
  - Clear boundaries on what can and cannot be done
  - Creating a solution within boundaries is a challenge
- MapReduce Design Patterns
  - A template for solving a common and general data manipulation problem with MapReduce
  - Is not specific to a domain, but a general approach to solving a problem

# BASIC PATTERNS

- Summarization
- Filtering
- Data Organization
- Join
- Metapatterns

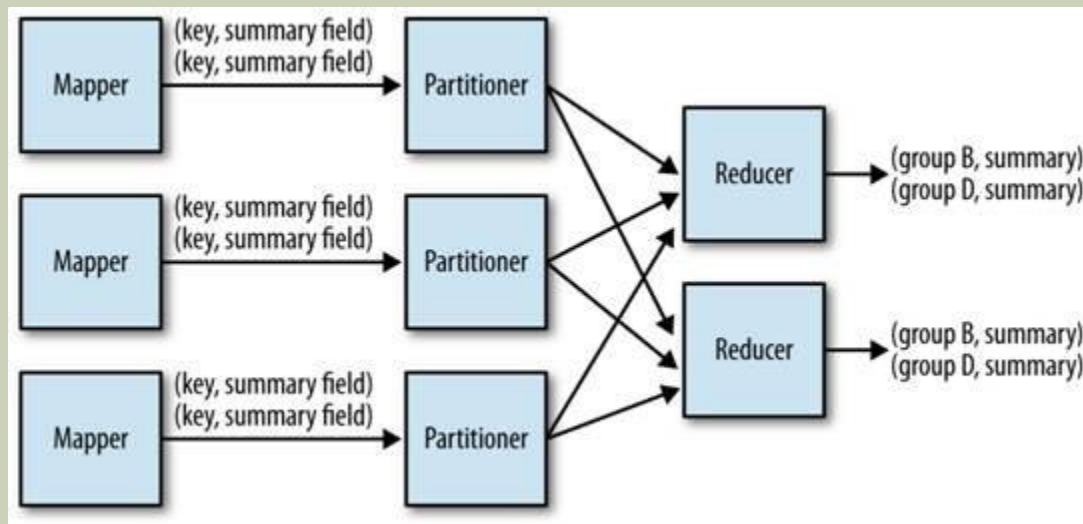
# SUMMARIZATION

- Grouping similar data together and performance an operation such as calculating a statistic, building an index, or just simply counting
- Examples
  - Numerical Summarizations
  - Inverted Index
  - Counting with Counters

# SUMMARIZATION

## NUMERICAL SUMMARIZATIONS

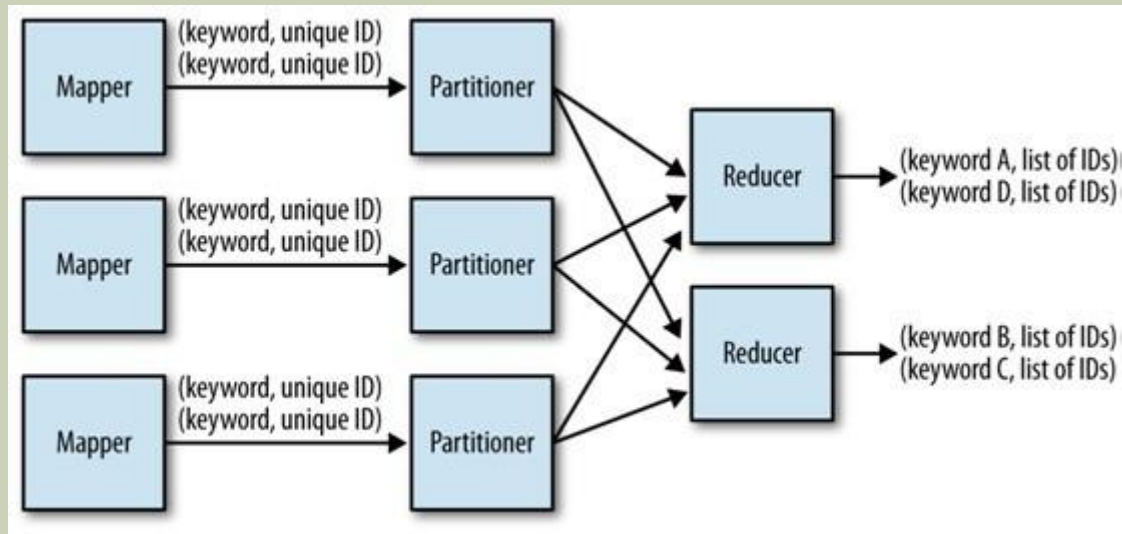
- Calculate aggregate statistical values over a large data set
- Benefit from a properly-designed combiner (How?)
- Benefit from a properly-designed partitioner (How?)



# SUMMARIZATION

## INVERTED INDEX

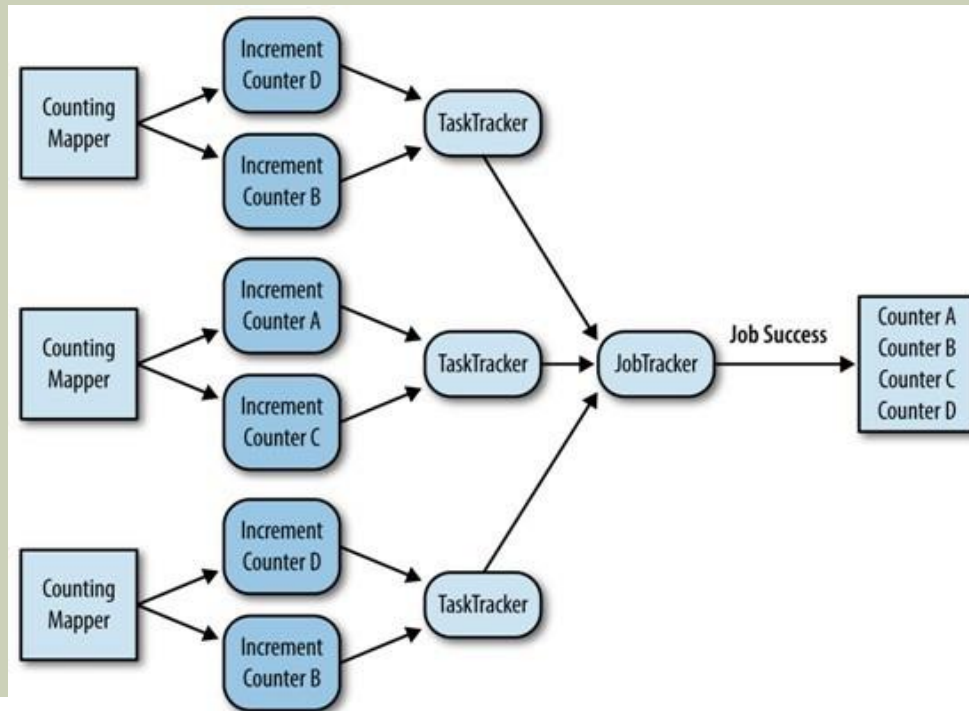
- Generating an index from a data set to allow for faster searches or data enrichment capabilities
- Benefit from the combiner (How?)



# SUMMARIZATION

## COUNTING WITH COUNTERS

- Utilizing MapReduce's internal counter to calculate global sums on the entire data set
- Done entirely in the mapping phase
- Useful if number of keys is small





# SUMMARIZATION

## COUNTING WITH COUNTERS

```
public static class CountNumUsersByStateMapper extends Mapper<Object, Text, NullWritable, NullWritable> {  
    public static final String STATE_COUNTER_GROUP = "State";  
    public static final String UNKNOWN_COUNTER = "Unknown";  
    public static final String NULL_OR_EMPTY_COUNTER = "Null or Empty";  
    private String[] statesArray = new String[] { "AL", "AK", ... };  
    private HashSet<String> states = new HashSet<String>(Arrays.asList(statesArray));  
  
    public void map(Object key, Text value, Context context) throws IOException, InterruptedException {  
        // Parsed data from value, get the value for the Location attribute, look for a state abbreviation  
        // code if the location is not null or empty  
        // For each token, check if it is a state. If so, increment the state's counter by 1 and flag it as  
        // not unknown  
        context.getCounter(STATE_COUNTER_GROUP, state).increment(1);  
        ...  
        // If the state is unknown, increment the UNKNOWN_COUNTER counter  
        context.getCounter(STATE_COUNTER_GROUP, UNKNOWN_COUNTER).increment(1);  
  
        // If it is empty or null, increment the NULL_OR_EMPTY_COUNTER counter by 1  
        context.getCounter(STATE_COUNTER_GROUP, NULL_OR_EMPTY_COUNTER).increment(1);  
    }  
}
```

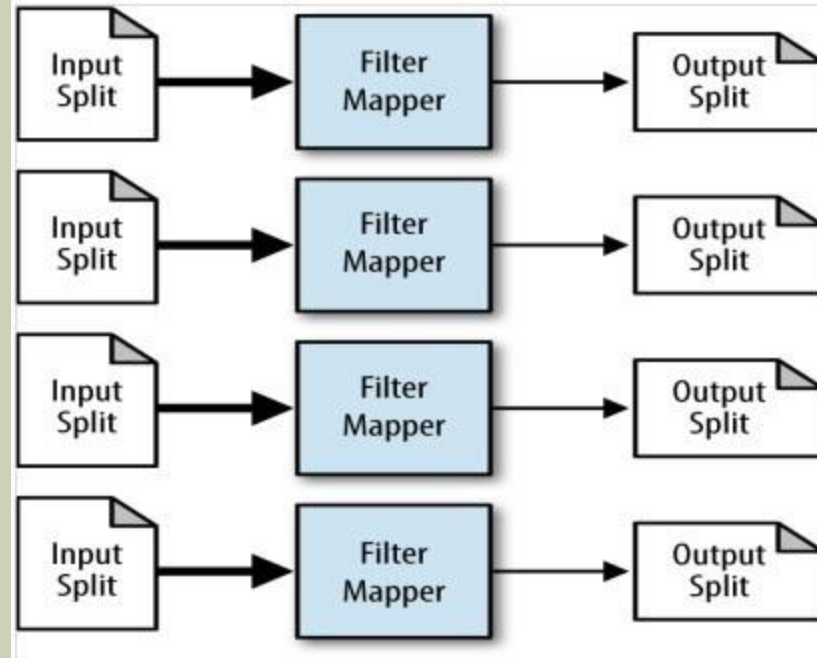
# FILTERING

- Don't change the actual records but find a subset of data in the entire data set
- Examples
  - Filtering
  - Top-ten
  - Distinct

# FILTERING

## FILTERING

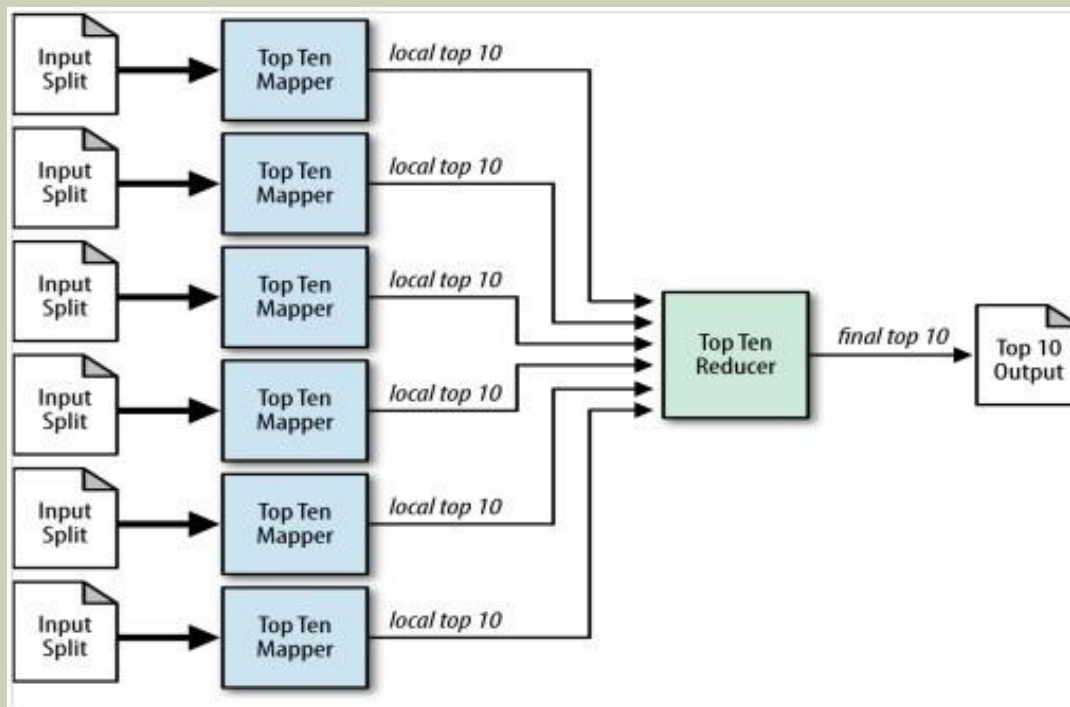
- Evaluates each record and decides, based on some condition, whether it should stay or go
- Not require the “reduce part”
- Can a combiner helps?



# FILTERING

## TOP-TEN

- Retrieve a relatively small number of top K records according to a ranking scheme in the data set
- Important: Unique entries



# **FILTERING**

## **DISTINCT**

- Find a unique set of values
- Key process: deduplication
- Benefit from deduplication processes implemented inside mappers and combiners
- Benefit from a large number of reducers

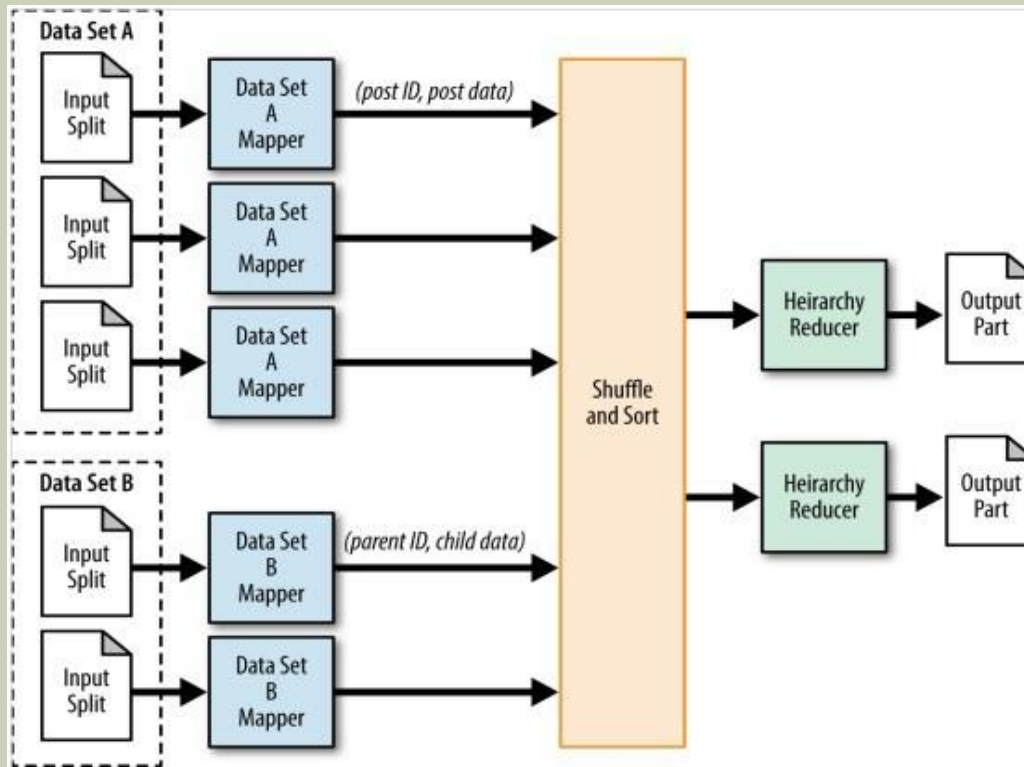
# DATA ORGANIZATION

- Reorganizing data to help improving value, performance, and usability.
- Examples:
  - Structured to Hierarchical
  - Partitioning and Binning

# DATA ORGANIZATION

## STRUCTURED TO HIERARCHICAL

- Create new records with a different structure from the input data (e.g.: transform row-based data into a hierarchical format such as XML or JSON)



# DATA ORGANIZATION

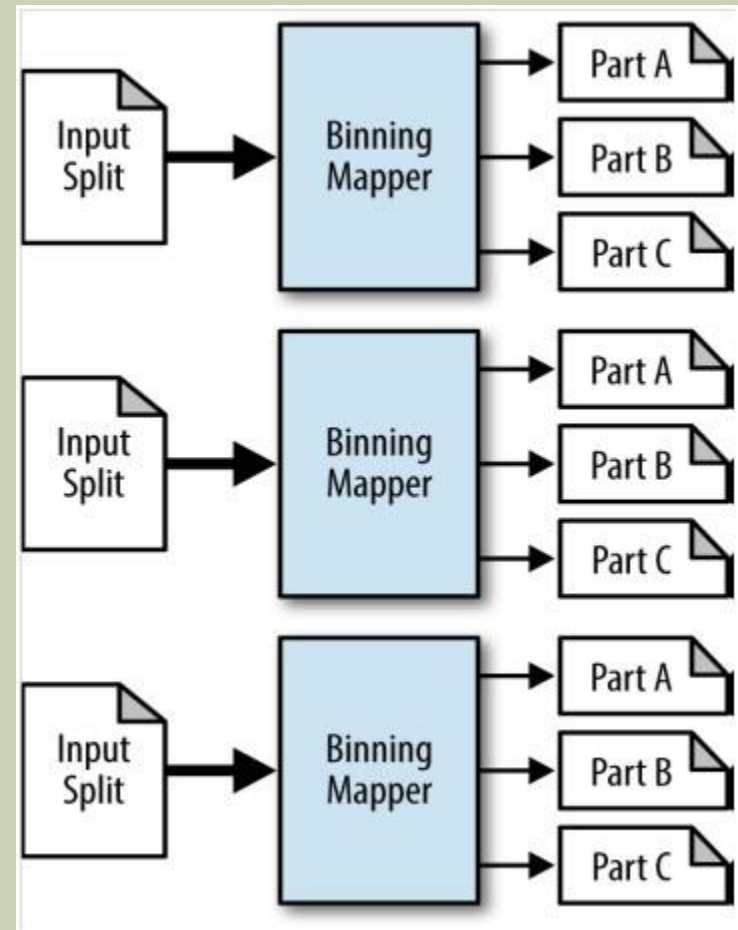
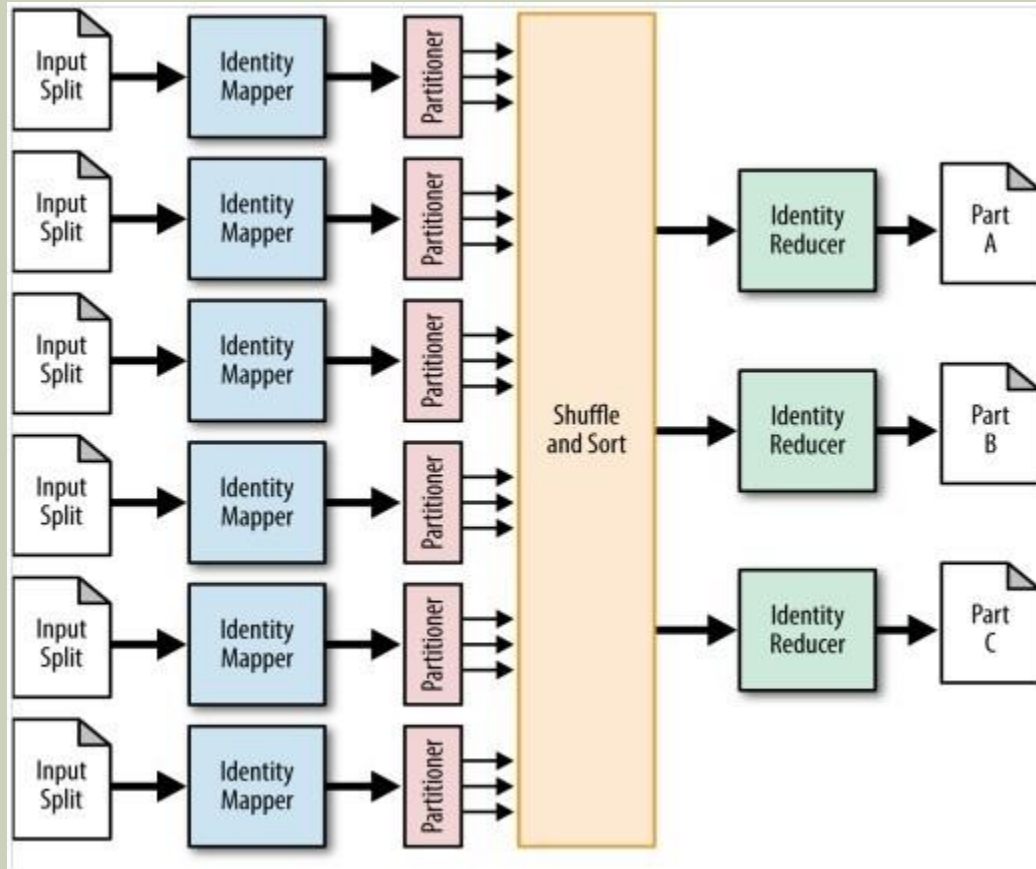
## PARTITIONING AND BINNING

- Relies on Hadoop's Partitioner and Binning capabilities to categorize data based on key ranges.
- Binning: Mapping phase
- Partitioning: After mapping



# DATA ORGANIZATION

## PARTITIONING AND BINNING



# JOIN

- Merging similar data (with or without aggregation) from multiple data sets
- Examples:
  - Reduce-side join
  - Map-side join
  - Cartesian product

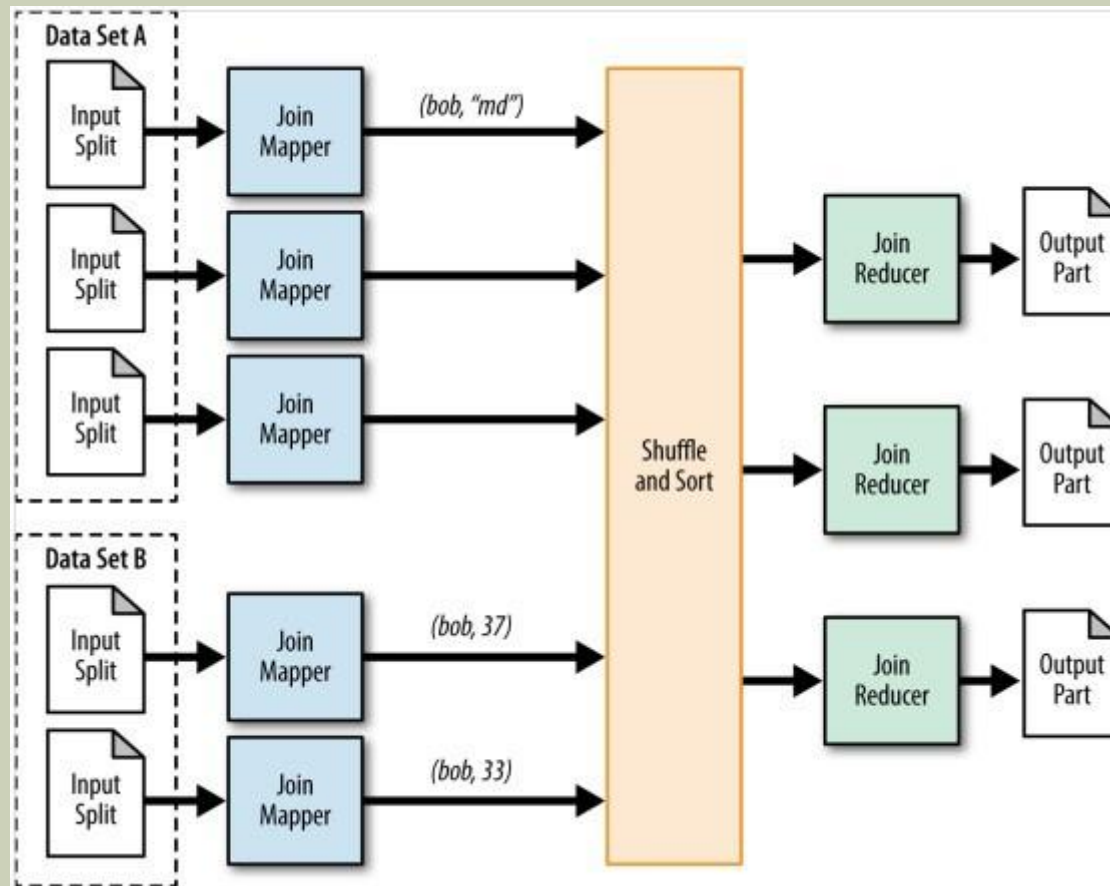
# JOIN

## REDUCE-SIDE JOIN

- Join large multiple data sets together by some foreign keys
- Simplest and most straight forward
- Should be the last solution to look at

# JOIN

## REDUCE-SIDE JOIN



# JOIN

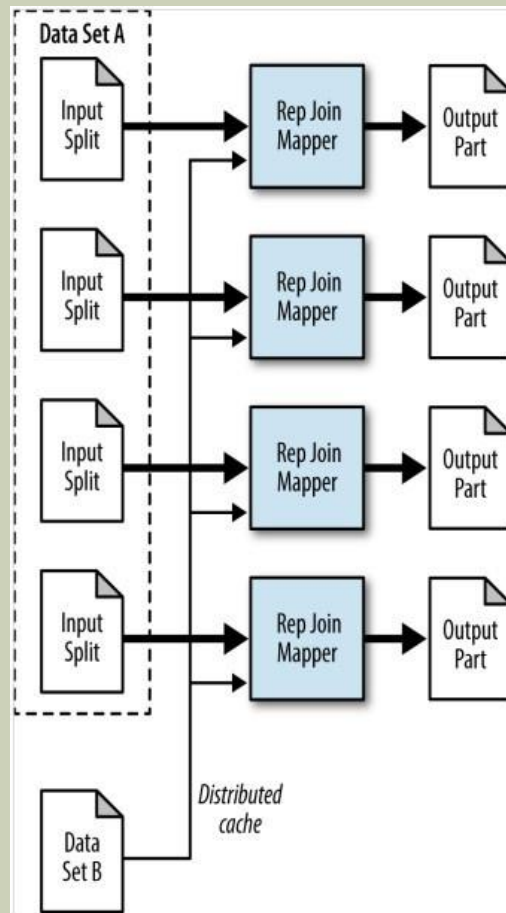
## MAP-SIDE JOIN

- Joining happens in the map phase
- Replicated Join
- Composite Join

# JOIN

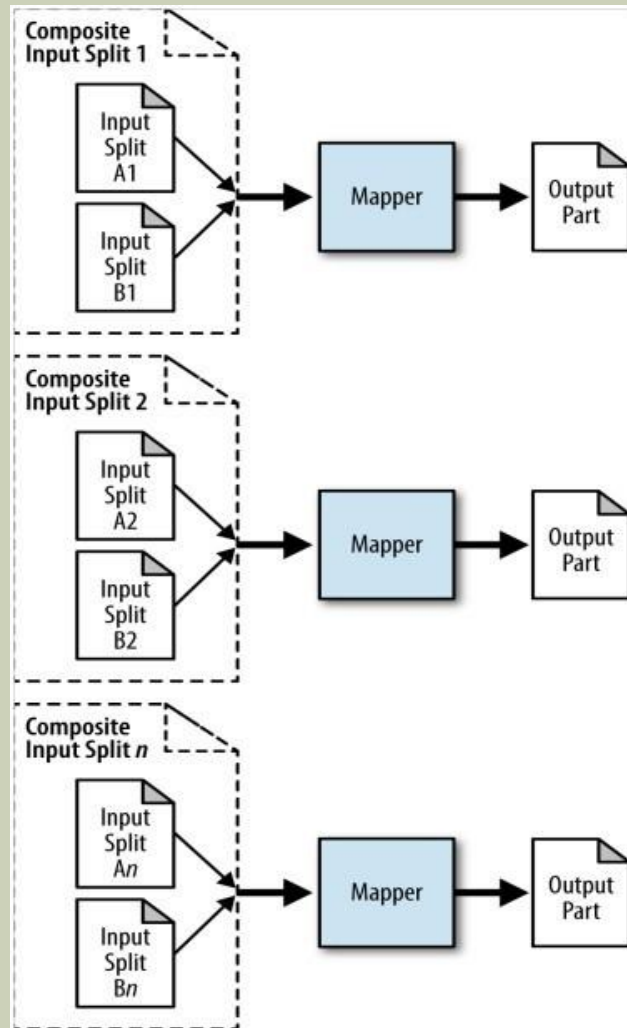
## MAP-SIDE JOIN: REPLICATED JOIN

- Use distributed cache on all mappers



# JOIN

## MAP-SIDE JOIN: COMPOSITE JOIN



# JOIN

## COMPOSITE JOIN

```
public static void main(String[] args) throws Exception {
    Path userPath = new Path(args[0]);
    Path commentPath = new Path(args[1]);
    Path outputDir = new Path(args[2]);
    String joinType = args[3];
    JobConf conf = new JobConf("CompositeJoin");
    conf.setJarByClass(CompositeJoinDriver.class);
    conf.setMapperClass(CompositeMapper.class);
    conf.setNumReduceTasks(0);
    // Set the input format class to a CompositeInputFormat class. The CompositeInputFormat will
parse all of our input files and output records to our mapper.
    conf.setInputFormat(CompositeInputFormat.class);
    // The composite input format join expression will set how the records are going to be read in,
and in what input format.
    conf.set("mapred.join.expr", CompositeInputFormat.compose(joinType,
        KeyValueTextInputFormat.class, userPath, commentPath));
    TextOutputFormat.setOutputPath(conf, outputDir);
    conf.setOutputKeyClass(Text.class);
    conf.setOutputValueClass(Text.class);
    RunningJob job = JobClient.runJob(conf);
}
```



# JOIN

## COMPOSITE JOIN

```
public static class CompositeMapper extends MapReduceBase  
implements
```

```
    Mapper<Text, TupleWritable, Text, Text> {
```

```
    public void map(Text key, TupleWritable value,  
        OutputCollector<Text, Text> output, Reporter reporter)  
        throws IOException {
```

```
        // Get the first two elements in the tuple and output them  
        output.collect((Text) value.get(0), (Text) value.get(1));
```

```
    }
```

```
}
```

# METAPATTERNS

- Patterns that deal with patterns
- Examples:
  - Job chaining
  - Job merging

# METAPATTERNS

## JOB CHAINING

- Many problems can't be solved with a single MapReduce job
- The default MapReduce framework requires a lot of manual coding to handle multistage jobs
  - Clean up intermediate outputs
  - Handle failures
- Can be chained through the main Java program or external scripting
- Possible support tools: Oozie, Cascade, Tez ...
- Serial chaining
- Parallel chaining

# METAPATTERNS

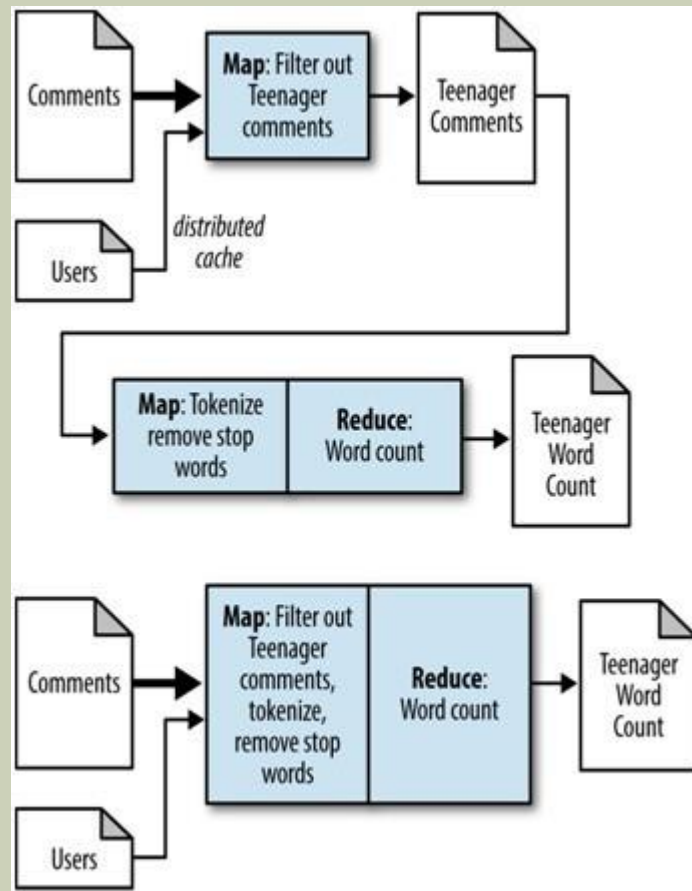
## JOB CHAINING: CHAIN FOLDING

- Optimization that is applied to MapReduce job chain
- Reduce the amount of data movement in the MapReduce pipe line

# METAPATTERNS

## JOB CHAINING: CHAIN FOLDING

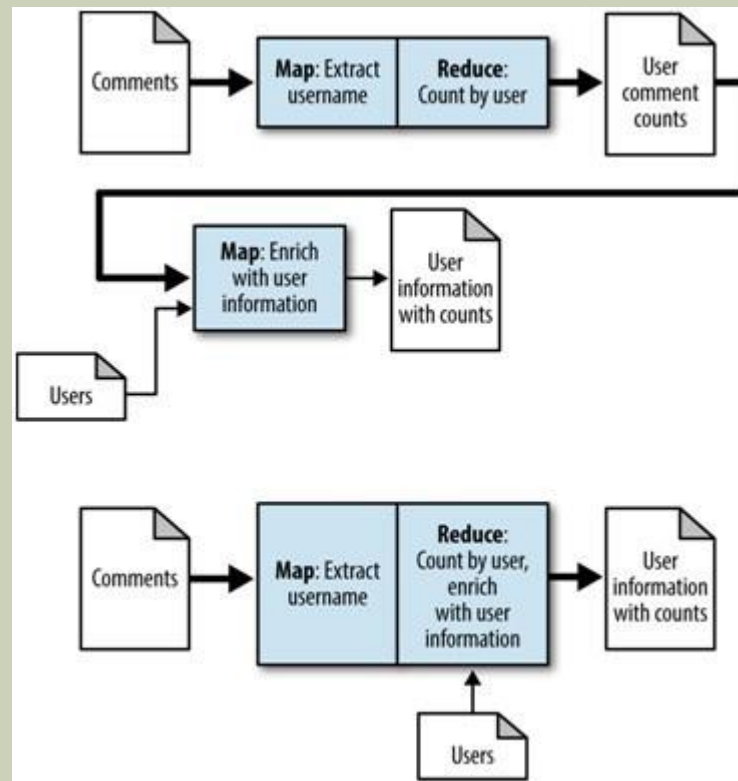
- If multiple map phases are adjacent, merge them



# METAPATTERNS

## JOB CHAINING: CHAIN FOLDING

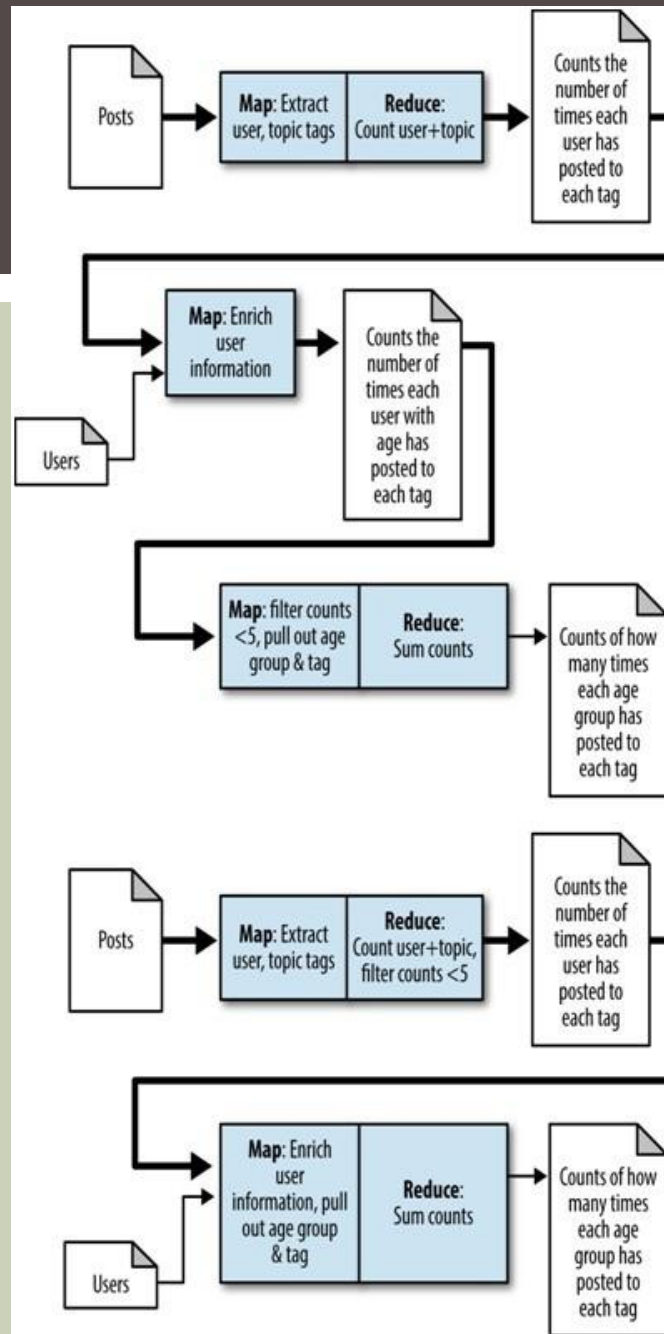
- If a job ends with a map phase, push that phase into the previous reducer



# METAPATTERNS

## JOB CHAINING: CHAIN FOLDING

- Split up map operations that *decrease* the amount of data from those that *increase* the amount of data





# METAPATTERNS

## JOB MERGING

- Allows unrelated jobs loading the same data to share the MapReduce pipeline

