



MapReduce Programming Framework



Basic Concepts

- What is “map”?
 - A function/procedure that is applied to every individual elements of a collection/list/array/...
 - `int square(x) { return x*x; }`
 - `map square [1,2,3,4] -> [1,4,9,16]`
- What is “reduce”?
 - A function/procedure that performs an operation on a list. This operation will “fold/reduce” this list into a single value (or a smaller subset)
 - `reduce ([1,2,3,4]) using sum -> 10`
 - `reduce ([1,2,3,4]) using multiply -> 24`



MPI: Perspective

- Call MPI_Scatter to distribute data
 - Knowledge about data size?
 - Bottleneck at root process
- Each process applies the “map” function on local data
 - + Processes can interact
- Call MPI_Reduce to applies the “reduce” function on final data
 - Need to define MPI_OP for complex reduce operation
 - Eventual bottleneck at root process for non-aggregative operations



Word Count

- Count how many unique words there are in a file/multiple files
- The “Hello World” of Big Data/Data Intensive Computing



Word Count: MPI

- MPI_Scatter:
 - Does data fit into memory?
 - If not, how do you set up MPI-IO or divide the files among the processes?
- Individual processes:
 - Unknown number of unique words in each workload: Linked list
 - `struct {char* word, int count}`
- MPI_Reduce
 - What happens if there are a lot of unique words?




Basic Programming Paradigm

- Programmers implement:
 - Map function:
 - Take in the input data and return a $\langle \text{key}, \text{value} \rangle$ pair
 - Reduce function:
 - Receive the $\langle \text{key}, \text{value} \rangle$ pairs from the mapper and provide a final output as a reduction operation on the pairs
- The MapReduce Framework (and HDFS) handles everything else



MapReduce WordCount Example

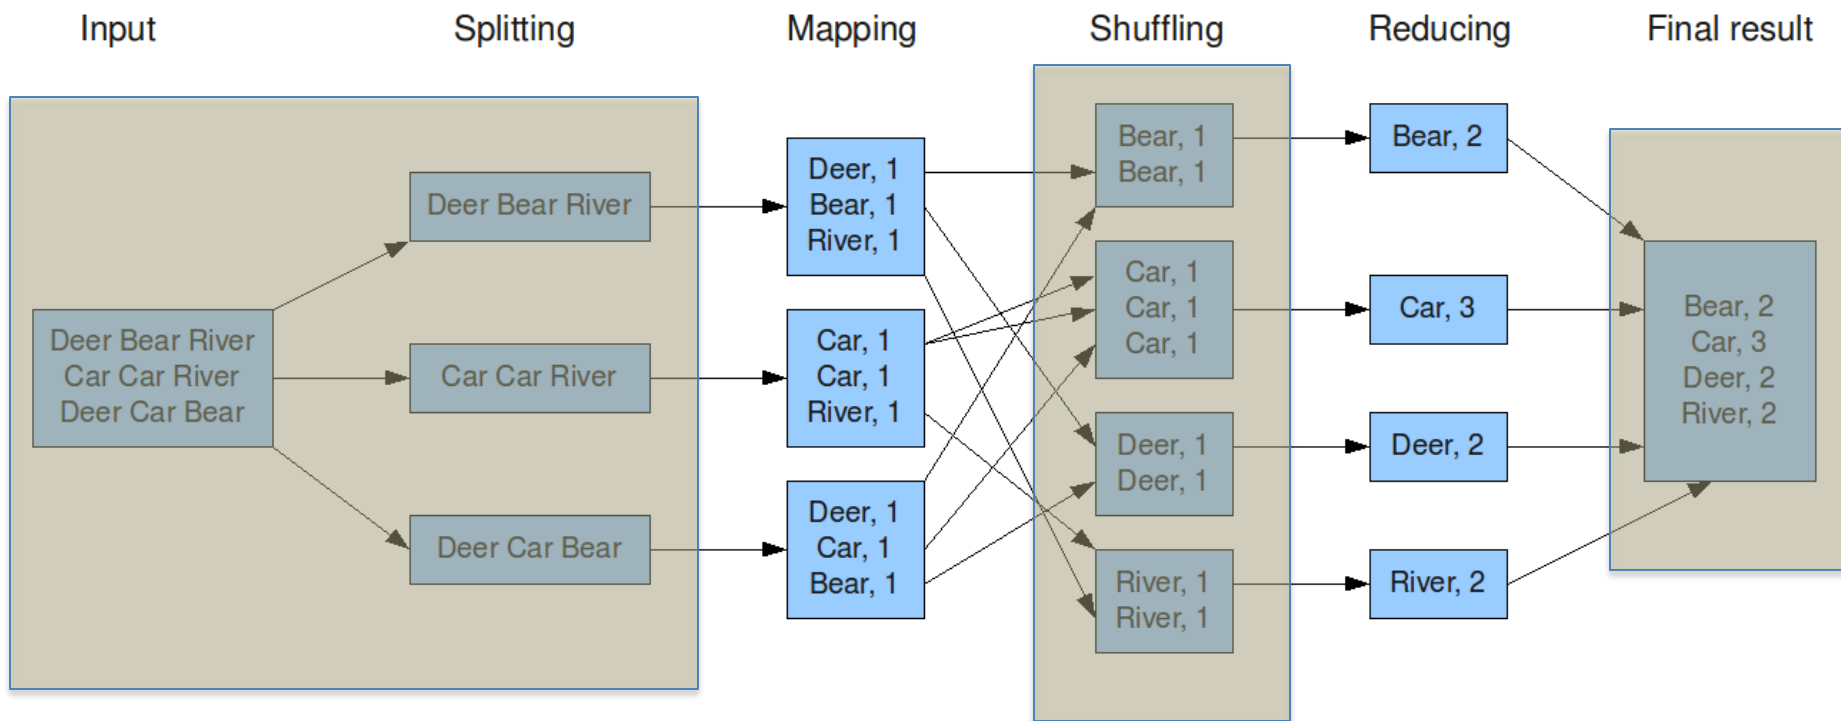
The overall MapReduce word count process





MapReduce WordCount Example

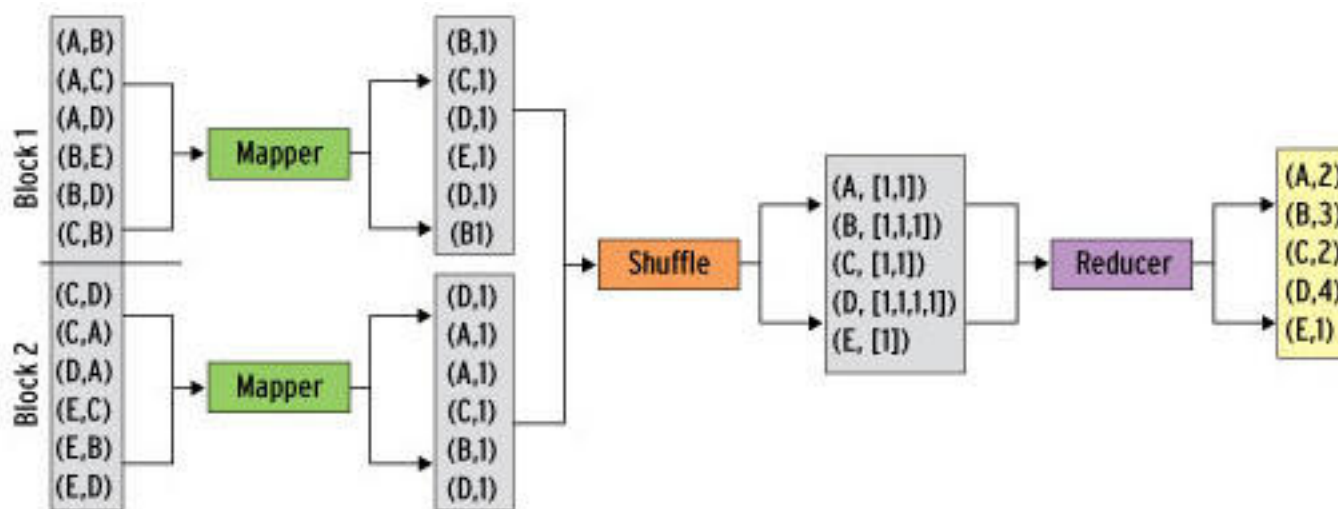
The overall MapReduce word count process

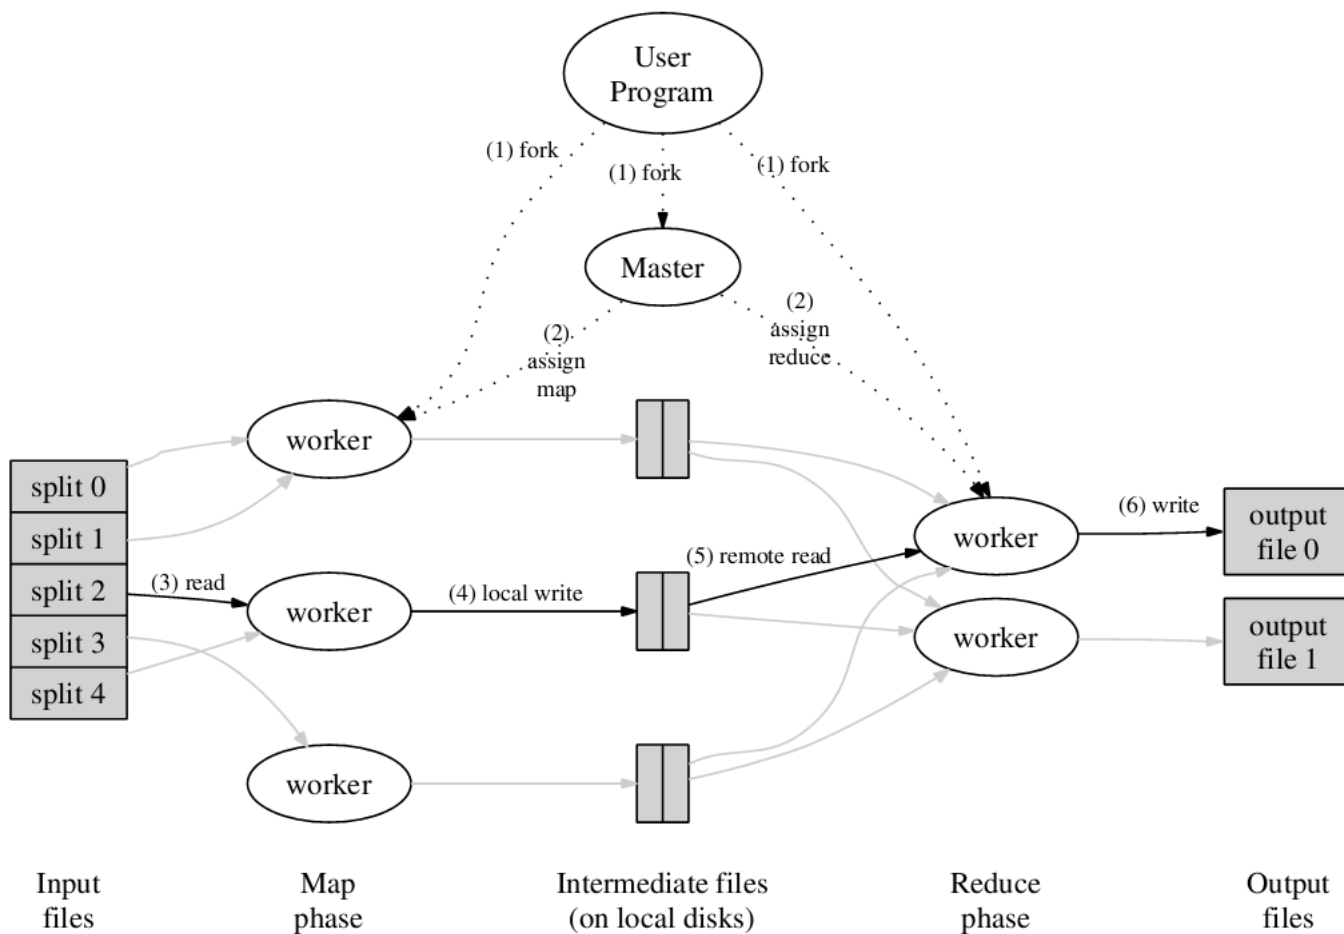


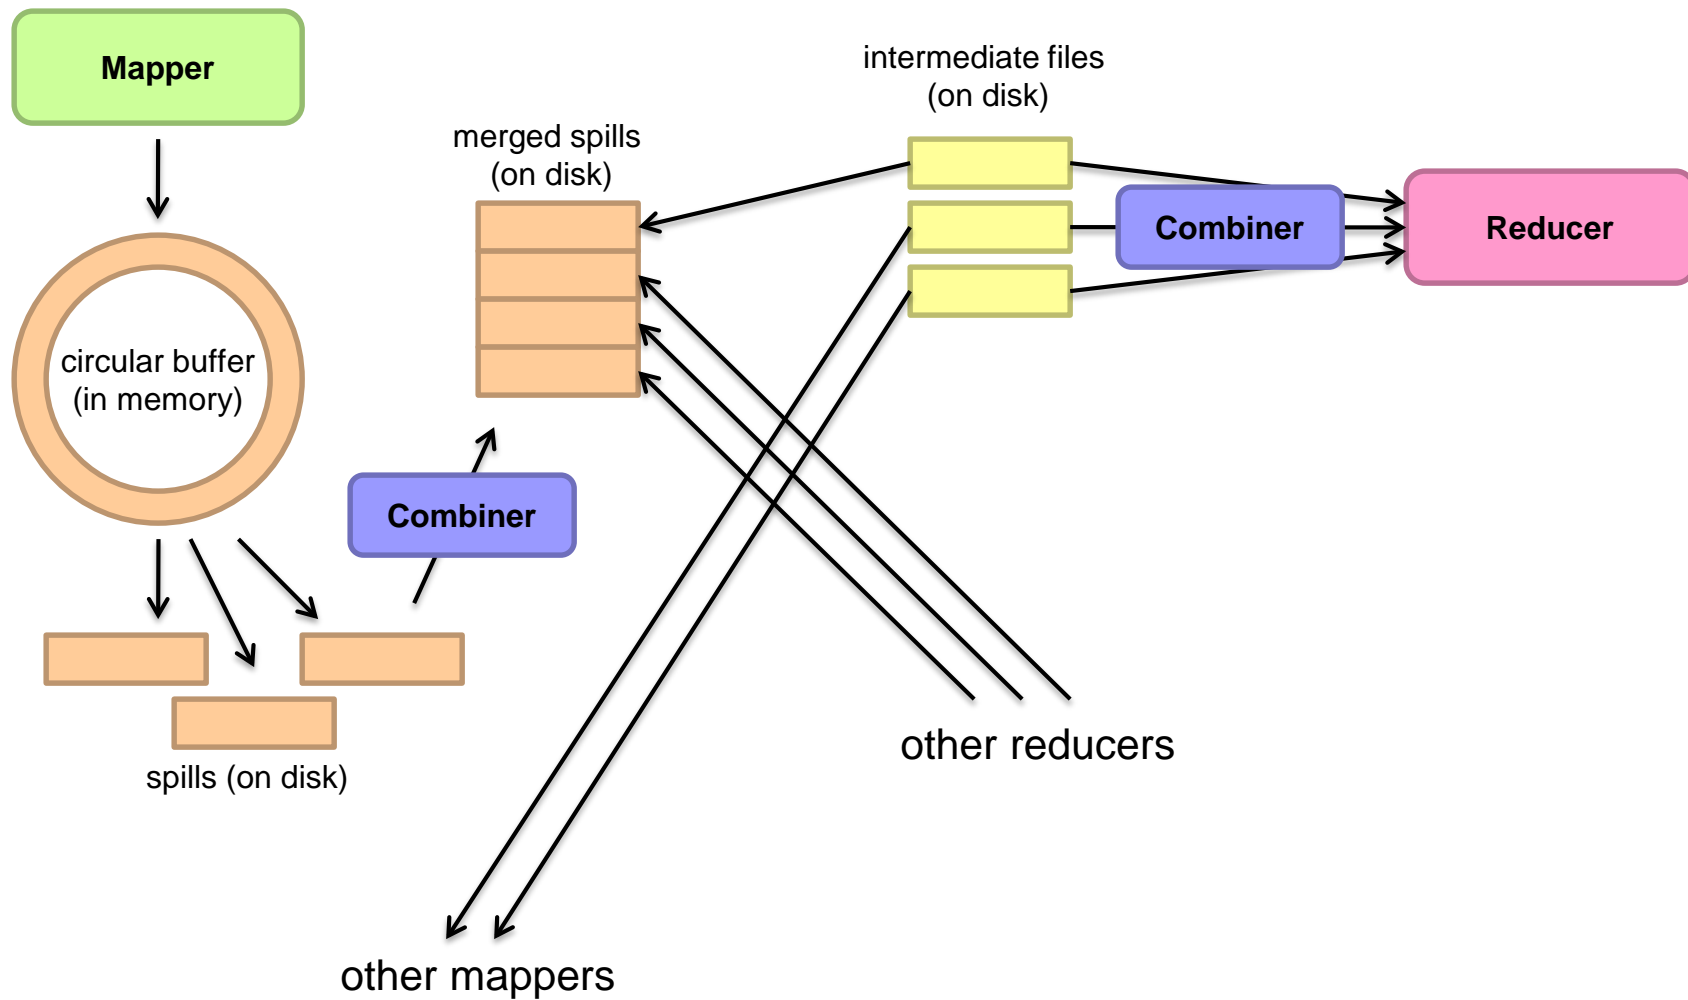


MapReduce PageRank Example 1

(A,B): There is a referral (link) from site A to site B. Google looks at how many referrals site B has in order to determine the ranking of site B.







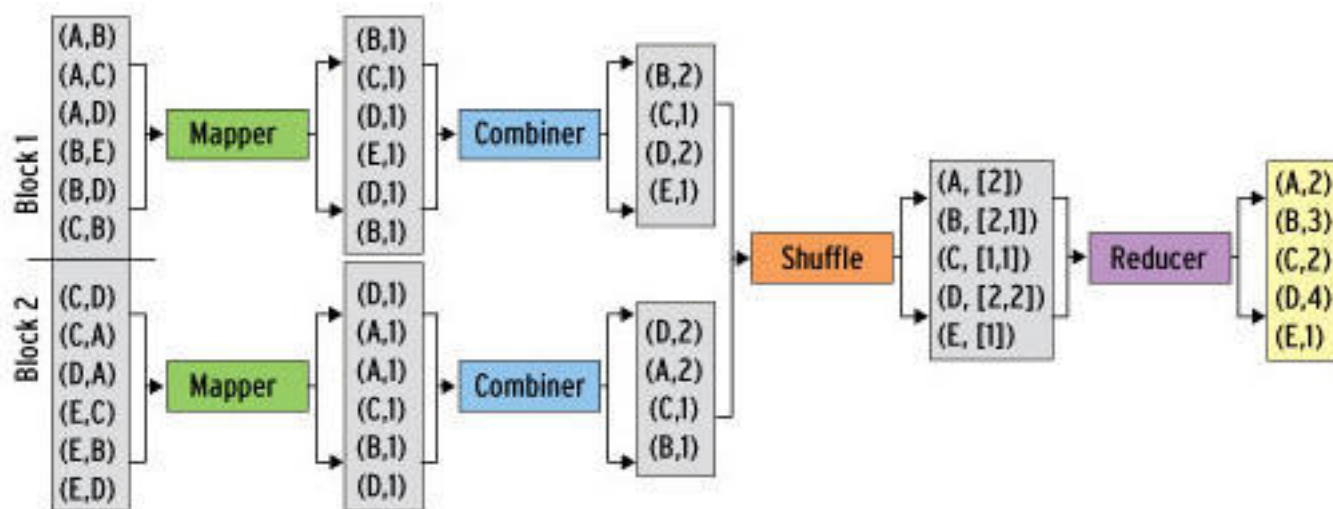


Basic Programming Paradigm

- Programmers implement:
 - Map function:
 - Take in the input data and return a <key,value> pair
 - Reduce function:
 - Receive the <key,value> pairs from the mapper and provide a final output as a reduction operation on the pairs
 - **Optional functions:**
 - **Partition function: determines the distribution of mappers' <key,value> pairs to the reducers**
 - **Combine functions: initial reduction on the mappers to reduce network traffics**
- The MapReduce Framework handles everything else



MapReduce PageRank Example 2





What is “everything else”?

- “Everything else”
 - Scheduling
 - Data distribution
 - Synchronization
 - Error and Fault Handling
- Limited control over data and execution flow
 - All algorithms must be expressed as a combination of mapping, reducing, combining, and partitioning functions
- Extremely limited knowledge on
 - Location of mappers and reducers
 - Life cycle of individual mappers and reducers
 - Information about which mapper handles which data block
 - Information about which reducer handles which intermediate key



Challenges in working with MR

- All algorithms must be expressed as a combination of mapping, reducing, and possibly combining and partitioning, functions
- Large scale debugging is difficult
 - Functional errors are difficult to follow at large scale
 - Data-dependent errors are even more difficult to catch and fix



Applications of MapReduce

- Text tokenization, indexing, and search
 - Web access log stats
 - Inverted index construction
 - Term-vector per host
 - Distributed grep/sort
- Graph creation
 - Web link-graph reversal
 - Google's PageRank
- Data Mining and machine learning
 - Document clustering
 - Machine learning
 - Statistical machine translation