


Memory-centric Distributed Computing / In-memory computing

Linh B. Ngo

MOTIVATION

Why do we want in-memory computing?

MOTIVATION



WHAT CAN MAPREDUCE OFFER?

■ Pros:

- Inherently parallel
- Data locality
- Resiliency and Error Recovery

■ Cons:

- High overhead costs for setting up (JVMs)
- Difficulty in setting up iterative jobs due to inability to transition between MR stages (having to write back to HDFS first)

DO WE NOT LIKE MAPREDUCE ANYMORE?

- “If all you have is a hammer, everything becomes a nail”
- MapReduce is good for massive scale batch-oriented, single iteration jobs (initial filtering, data cleaning, ...)
- Additional tools must be developed

WHAT HAVE BEEN DEVELOPED

- **Apache Tez**
 - Enable automated “chaining” of MR jobs via dataflow graphs
- **Apache Spark**
 - Fast processing via read-only in-memory data cache
- **Tachyon**
 - RAMdisk-based storage on top of persistent storage (HDFS) to support fast in-memory data access

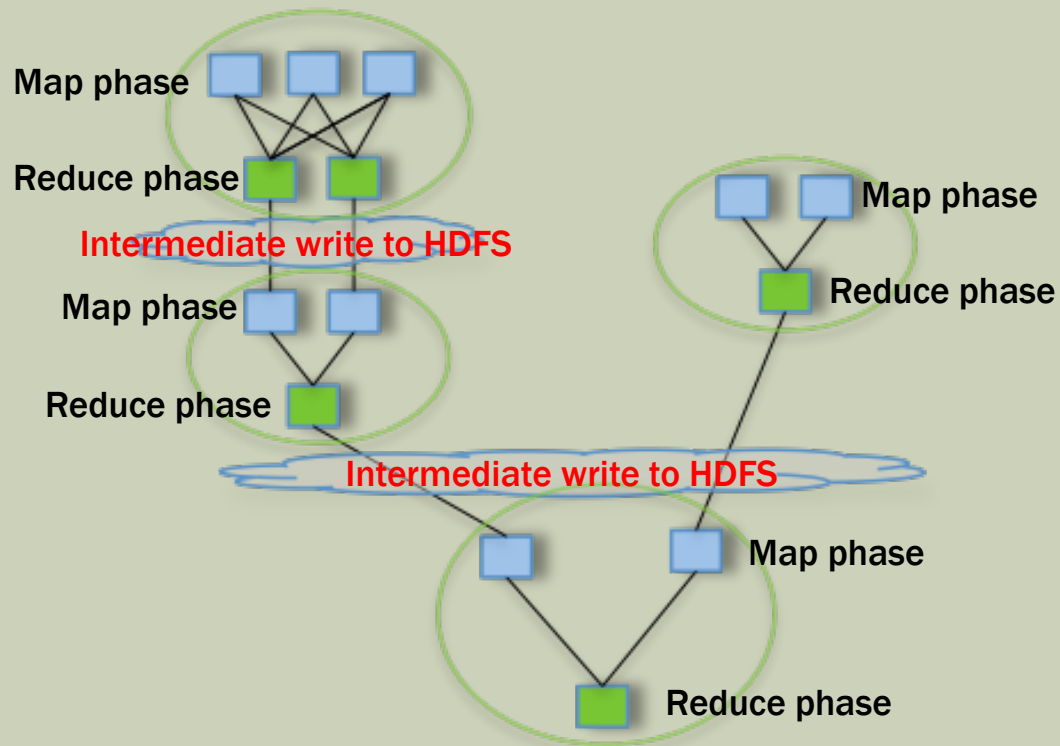
APACHE TEZ

- Extensible framework for building high performance batch and interactive data processing applications
- User-defined flows of data and tasks to support complex query logic that typically require multiple MR jobs
- Current status: Apache top-project, 0.8.4

APACHE TEZ

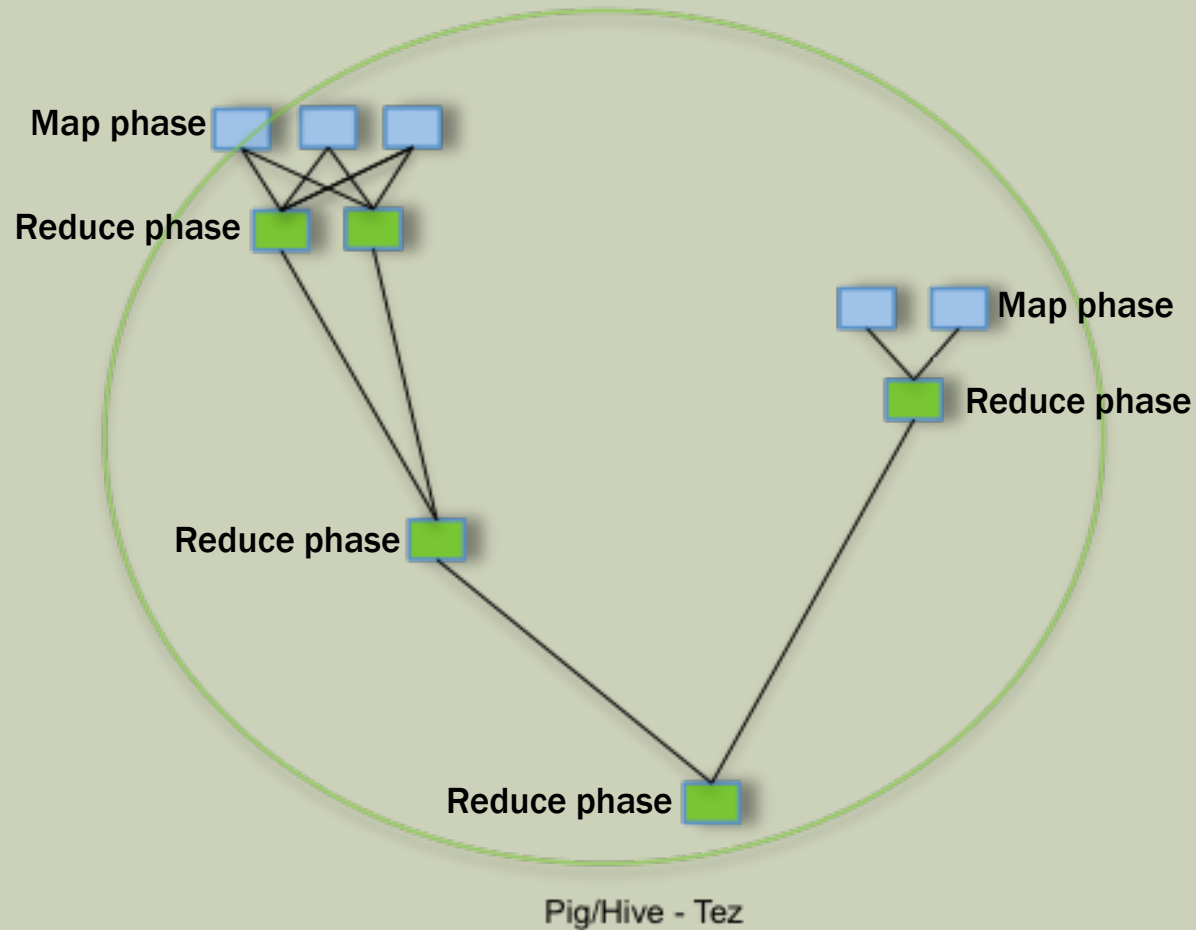
- Data processing is represented as a dataflow graph
 - Vertices represent the processing of data
 - Edges represent movement of data between the processing
- For MR jobs, Tez provides the ability to run a chain of Reduce stages as compared to a single Reduce stage

APACHE TEZ



Pig/Hive - MR

APACHE TEZ



APACHE TEZ: BUILDING BLOCKS

- **Vertex:** Combination of a set of Inputs, a Processor, and a set of Outputs
 - **Input:** A pipe through which a processor can accept input data from a data source
 - **Processor:** The entity responsible for consuming inputs and producing outputs
 - **Output:** A pipe through which a processor can generate output data

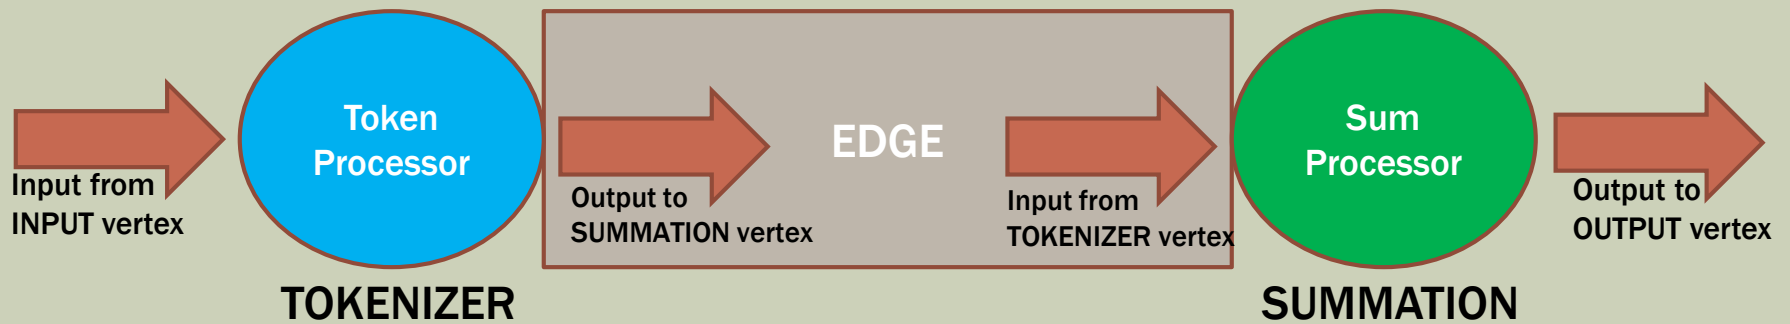
APACHE TEZ: BUILDING BLOCKS

- **Edge:** A logical entity that represents a number of physical connections between the tasks of two connected vertices
 - **Logical:** representing the logical edge between 2 vertices
 - **Physical:** representing the connection between a task of the Source vertex and a task of a Destination vertex
 - A Logical Edge can contain multiple Physical Connections

WORDCOUNT DAG

INPUT: default input from files

OUTPUT: default output to be written to file



APACHE SPARK

- “Lightning-fast cluster computing”
- Integrated into Hortonwork’s and Cloudera’s Hadoop 2.0 production release
- 1.6.0

TARGET SYSTEMS/APPLICATIONS

■ Systems:

- COTS Clusters for large scale data analytics
- Highly scalable

■ Applications:

- Reuse intermediate results across multiple computations
- Common in many iterative machine learning and graph algorithms

RESILIENT DISTRIBUTED DATASETS (RDD)

- An abstraction to enable efficient data reuse in a broad range of application
- RDD is implemented in Spark

CHARACTERISTICS OF RDD

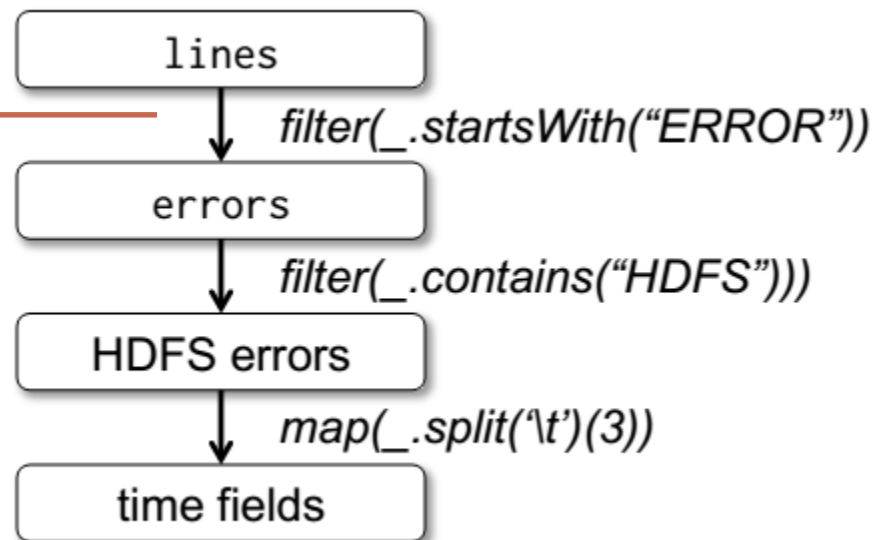
- Read-only, partitioned collection of records
- Created (aka *written*) through deterministic operations on data
 - In stable storage
 - From other RDDs
 - Coarse-grained operations: map, join, filter ...
- Do not need to be materialized at all time
 - RDDs are recoverable via **data lineage**

CHARACTERISTICS OF RDD

```
lines = spark.textFile("hdfs://...")
errors = lines.filter(_.startsWith("ERROR"))
errors.persist()
```

Stable data in
HDFS

Persistence in-
memory RDD



APPLICATIONS NOT SUITABLE FOR SPARK

- Applications that make asynchronous fine-grained updates to shared state:
 - Storage system for web applications (Traditional applications for standard databases)
 - Incremental web crawler
 - ...

APPLICATIONS SUITABLE FOR SPARK

- Batch Analytics/Pleasantly Parallel Applications
- Examples:
 - Logistic Regression
 - PageRank
 - Machine Learning

TACHYON

- Reliable data sharing at memory speed across cluster computing framework
- Integrated into Bekerley Data Analytics Stack (BDAS)
- 0.8

TACHYON: MOTIVATION

- Caching data in memory during computation
- Need in-memory data storage
- Similar to Spark's concepts:
 - Tachyon stays in memory
 - No data replication, only data lineage
 - Raw data are loaded into Tachyon from persistent storage on disk (e.g.: HDFS)

TACHYON: DESIGN ASSUMPTIONS

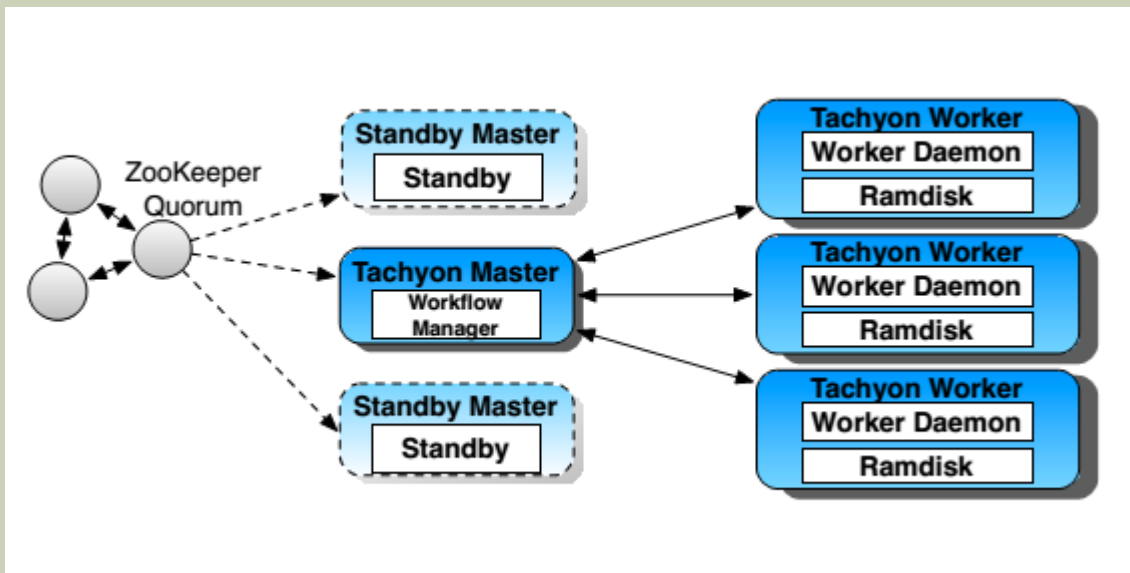
- Immutable data (write once, read many)
- Deterministic jobs (multiple runs produce same results)
- Locality-based scheduling
- Working set data fits in memory (Full data requires disk storage)
- Program size \ll data size (move computation to data)

TACHYON: KEY ARCHITECTURAL DETAILS

- Two layers
- Master-slave Architecture
- Lineage Overhead Garbage Collection
- Data Eviction
- Master Fault-Tolerance
- Check-pointing
- Resource Allocation

TACHYON: KEY ARCHITECTURAL DETAILS

- Two layers
 - Lineage Layer: In-memory Tachyon daemons
 - Persistent Layer: Replication-based storage systems (HDFS, S3, ...)
- Master-slave Architecture
 - Master: Lineage metadata and workflow manager
 - Workers: memory-mapped files in RAMdisk



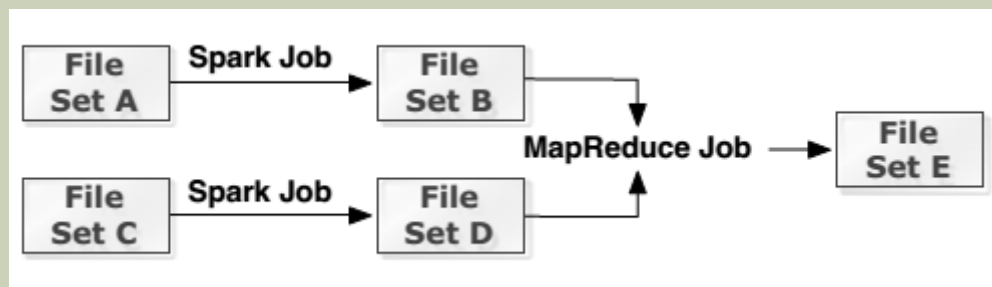
TACHYON: KEY ARCHITECTURAL DETAILS

■ Lineage Overhead:

- Metadata
- Job binaries
- Garbage collected

■ Data Eviction

- Access frequency
- Access temporal locality



Return	Signature
Global Unique Lineage Id	<code>createDependency(inputFiles, outputFiles, binaryPrograms, executionConfiguration, dependencyType)</code>
Dependency Info	<code>getDependency(lineageId)</code>

TACHYON: KEY ARCHITECTURAL DETAILS

■ Master Fault-Tolerance

- Standby secondary master
- Zookeeper to elect new master
- Persistent layer to rebuild from lineage

■ Check-pointing

- Check-point for the entire Tachyon system, not just application
- Check-point key elements in the lineage chain
- Check-point hot files
- Avoid check-point temporary files

TACHYON: KEY ARCHITECTURAL DETAILS

■ Resource Allocation

- How to balance resources between actual work and recomputation work
- Priority compatibility: recomputation priority follow job priority
- Resource sharing (no fixed allocation to recomputation)
- Avoid cascading recomputation

TACHYON: LIMITATION/FUTURE WORK

- Enabling random access abstraction
- Enabling mutable data
- Enabling multi-tenancy (memory sharing among users)
- Hierarchical storage (NVRAM, SSDs ...)
- Optimizing checkpoint algorithm

FUTURE DIRECTION

- Discardable Distributed Memory: Supporting Memory Storage in HDFS
- Taking lessons from Spark and Tachyon (Hortonworks)
- To be integrated into HDFS core libraries